

Documentation for Phase3 Operating System Project - A.Y. 2023/2024

Introduction

The following document serves as a simple explanation of the main features of the third phase of the OS project for the *uRISCv* architecture.

Index

- [Virtual Memory Support](#)
- [User Exception Handler](#)
- [System Service Thread](#)
- [Stdlib](#)
- [P3test](#)

Virtual Memory Support

Pager

The pager is a module that is responsible for managing the virtual memory of the system. It is responsible for loading and unloading pages from the disk to the memory, using the swap pool as a support struct for handling pages.

```

void pager(void) {
    unsigned status;
    // get the support data of the current process
    support_t *support_data = getSupportData();

    state_t *exception_state = &(support_data->sup_exceptState[PGFAULTEXCEPT]);

    unsigned cause = exception_state->cause & 0x7FFFFFFF;

    // check if the exception is a TLB-Modification exception
    if (cause == TLBMOD) {
        // treat this exception as a program trap
        programTrapExceptionHandler(exception_state);
    }

    gainSwapMutex();

    /* enter the critical section */

    // it's not the actual vpn, but the page index in the backing store
    int vpn = ENTRYHI_GET_VPN(exception_state->entry_hi);
    if (vpn >= MAXPAGES) {
        vpn = MAXPAGES - 1;
    }

    // pick a frame from the swap pool
    unsigned victim_frame = getFrameFromSwapPool();
    memaddr victim_page_addr = SWAPPOOLADDR + (victim_frame * PAGESIZE);

    // check if the frame is occupied
    if (isSwapPoolFrameFree(victim_frame) == FALSE) {
        // we assume the frame is occupied by a dirty page

        // operations performed atomically
        OFFINTERRUPTS();

        // mark the page pointed by the swap pool as not valid
        swap_pool[victim_frame].sw_pte->pte_entryLO &= !VALIDON;

        // update the TLB if needed
        updateTLB(swap_pool[victim_frame].sw_pte);

        ONINTERRUPTS();

        // update the backing store
        status = writeBackingStore(victim_page_addr, swap_pool[victim_frame].sw_asid,
                                   swap_pool[victim_frame].sw_pageNo);
        if (status != DEVRDY) {
            programTrapExceptionHandler(exception_state);
        }
    }

    // read the contents of the current process's backing store
    status =
        readBackingStoreFromPage(victim_page_addr, support_data->sup_asid, vpn);

    if (status != DEVRDY) // operation failed
    {

```

```

    programTrapExceptionHandler(exception_state);
}

// update the swap pool table
swap_pool[victim_frame].sw_asid = support_data->sup_asid;
swap_pool[victim_frame].sw_pageNo = vpn;
swap_pool[victim_frame].sw_pte = &support_data->sup_privatePgTbl[vpn];

// #region atomic operations
OFFINTERRUPTS();

// update the current process's page table
support_data->sup_privatePgTbl[vpn].pte_entryLO |= DIRTYON;
support_data->sup_privatePgTbl[vpn].pte_entryLO |= VALIDON;
support_data->sup_privatePgTbl[vpn].pte_entryLO &= 0xfff;
support_data->sup_privatePgTbl[vpn].pte_entryLO |= victim_page_addr & 0xfffff000;

// update the TLB
updateTLB(&support_data->sup_privatePgTbl[vpn]);

ONINTERRUPTS();
// #endregion atomic operations

releaseSwapMutex();
/* exit the critical section */

// return control to the current process
LDST(exception_state);
}

```

The pager is responsible for handling page faults. When a page fault occurs, the pager is invoked and it is responsible for loading the page from the backing store to the memory. The pager is also responsible for updating the TLB and the page table of the current process. It's important to ensure that the operations performed by the pager are atomic, so we disable the interrupts before performing the operations and re-enable them after the operations are completed (see the `OFFINTERRUPTS` and `ONINTERRUPTS` macros).

```

#define OFFINTERRUPTS() setSTATUS(getSTATUS() & (~MSTATUS_MIE_MASK))
#define ONINTERRUPTS() setSTATUS(getSTATUS() | MSTATUS_MIE_MASK)

```

We also need to ensure that the operations performed by the pager are thread-safe, so we use a mutex to ensure that only one pager can run at a time: that's accomplished by calling the `gainSwapMutex` and `releaseSwapMutex` functions, which acquire and release the mutex using message passing.

Swap Pool

The swap pool is a data structure that is used to store the pages that are swapped out of the memory. It is used by the pager to store the pages that are evicted from the memory when a page fault occurs. The swap pool is implemented as an array of swap pool entries, where each entry contains the address of the page in the backing store, the ASID of the process that owns the page, and the page number of the page in the backing store.

```

typedef struct swap_t
{
    int sw_asid;          /* ASID number */
    int sw_pageNo;        /* page's virt page no. */
    pteEntry_t *sw_pte;   /* page's PTE entry. */
} swap_t;

```

Page Tables

Are part of the process control block (support struct) and are used to store the mapping between the virtual pages and the physical frames. The page table is implemented as an array of page table entries, where each entry contains the physical frame number, the ASID of the process that owns the page and the process register `entryHI` and `entryLO`.

```
typedef struct pteEntry_t
{
    unsigned int pte_entryHI;
    unsigned int pte_entryLO;
} pteEntry_t;
```

Page Replacement Algorithms

The pager uses a page replacement algorithm to decide which page to evict from the memory when a page fault occurs. The algorithm looks for a free frame in the swap pool and if there are no free frames, it selects one in a round-robin fashion. The algorithm is implemented in the `getFrameFromSwapPool` function.

```
unsigned getFrameFromSwapPool() {
    static unsigned frame = 0;
    // find a free frame in the swap pool
    for (unsigned i = 0; i < POOLSIZE; i++) {
        if (isSwapPoolFrameFree(i)) {
            frame = i;
            break;
        }
    }
    // otherwise implement the page replacement algorithm RR
    return frame++ % POOLSIZE;
}
```

Read/Write to Backing Store

We use the `readBackingStoreFromPage` and `writeBackingStore` functions to read and write to the backing store. The backing store is a disk that is used to store the pages that are swapped out of the memory. The backing store is implemented as a file that is stored on the disk. We use the ssi apis to read and write to the backing store using the mnemonics `DOIO` operation. Is important to note which how this operation is performed:

1. We set the page address in the flash device register
2. We set the command value in the flash device register
3. We send a message to the ssi device to perform the operation
4. We receive the status of the operation

```
unsigned flashOperation(unsigned command, unsigned page_addr, unsigned asid,
                        unsigned page_number) {
    dtpreg_t *flash_dev_addr = (dtpreg_t *)DEV_REG_ADDR(IL_FLASH, asid - 1);
    flash_dev_addr->data0 = page_addr;

    unsigned value = (page_number << 8) | command;
    unsigned status = 0;
    ssi_do_io_t do_io = {
        .commandAddr = &(flash_dev_addr->command),
        .commandValue = value,
    };
    ssi_payload_t payload = {
        .service_code = DOIO,
        .arg = &do_io,
    };
    SYSCALL(SENDMESSAGE, (unsigned int)ssi_pcb, (unsigned int)&payload, 0);
    SYSCALL(RECEIVEMESSAGE, (unsigned int)ssi_pcb, (unsigned int)&status, 0);
    return status;
}
```

The two functions:

- `readBackingStoreFromPage`
- `writeBackingStore`

act as a wrapper around the `flashOperation` function, they are used to read and write to the backing store.

TLB Update

This is one of the most important parts of the pager, the TLB is a cache that is used to store the mappings between the virtual pages and the physical frames. As the TLB is implemented as an array of TLB entries, we can access it using built in macros. The *TLB update* is by performing a check if the page is already in the TLB, if it is, we update the entry, otherwise we add a new entry to the TLB.

```
void updateTLB(pteEntry_t *page) {
    // place the new page in the Data0 register
    setENTRYHI(page->pte_entryHI);
    TLBP();
    // check if the page is already in the TLB
    unsigned is_present = getINDEX() & PRESENTFLAG;
    if (is_present == FALSE) {
        // the page is not in the TLB, so we need to insert it
        setENTRYHI(page->pte_entryHI);
        setENTRYLO(page->pte_entryLO);
        TLBWI();
    }
}
```

User Exception Handler

The user exception handler is a module that is responsible for handling exceptions that occur in user mode. It is responsible for handling exceptions such as system calls, it is implemented as a switch statement that checks the cause of the exception and calls the appropriate handler function.

```
...
switch (exception_code)
{
case SYSEXCEPTION:
    UsysCallHandler(exception_state, current_support->sup_asid);
    break;
default:
    programTrapExceptionHandler(exception_state);
    return;
    break;
}
...
```

User System Call Handler

The user system call handler act as a wrapper for communication between user process and kernel, in fact it 'wraps' both the `SENDMESSAGE`

```

...
case SENDMSG:
    /* This services cause the transmission of a message to a specified process.
     * The USYS1 service is essentially a user-mode "wrapper" for the
     * kernel-mode restricted SYS1 service. The USYS1 service is requested by
     * the calling process by placing the value 1 in a0, the destination process
     * PCB address or SST in a1, the payload of the message in a2 and then
     * executing the SYSCALL instruction. If a1 contains PARENT, then the
     * requesting process send the message to its SST [Section 6], that is its
     * parent.
     */

    dest_process =
        a1_reg == PARENT ? sst_pcb[asid-1] : (pcb_t *)a1_reg;

    SYSCALL(SENDMESSAGE, (unsigned)dest_process, a2_reg, 0);

    break;
...

```

and the RECEIVEMESSAGE system calls.

```

...
case RECEIVMSG:
    /* This system call is used by a process to extract a message from its inbox
     * or, if this one is empty, to wait for a message. The USYS2 service is
     * essentially a user-mode "wrapper" for the kernel-mode restricted SYS2
     * service. The USYS2 service is requested by the calling process by placing
     * the value 2 in a0, the sender process PCB address or ANYMESSAGE in a1, a
     * pointer to an area where the nucleus will store the payload of the
     * message in a2 (NULL if the payload should be ignored) and then executing
     * the SYSCALL instruction. If a1 contains a ANYMESSAGE pointer, then the
     * requesting process is looking for the first message in its inbox, without
     * any restriction about the sender. In this case it will be frozen only if
     * the queue is empty, and the first message sent to it will wake up it and
     * put it in the Ready Queue.
     */

    receive_process = a1_reg == PARENT ? sst_pcb[asid-1] : (pcb_t *)a1_reg;
    SYSCALL(RECEIVEMESSAGE, (unsigned)receive_process, a2_reg, 0);

    break;
...

```

System Service Thread

The System Service Thread (SST) is a per-process thread that provide is child process useful services. Each SST child process can send a message to its SST to request a service (that can then be asked to the SSI if needed). The SST also initialize its child and Each share the same ID (ASID) and support struct of its child U-proc. Like the SSI the structure of SST works as a server: get the request, satisfy request and send back resoult. The SST is initialized by the `sstEntry` function, that initializes the child process, the print process and the term process. It then enters a loop where it waits for messages from the child process, handles the request and sends back the result.

```

void sstEntry() {
    // init the child
    support_t *sst_support = getSupportData();
    child_pcb[sst_support->sup_asid - 1] = initUProc(&u_proc_state[sst_support->sup_asid - 1], sst_support);

    // init the print process
    print_pcb[sst_support->sup_asid - 1] = initPrintProcess(&print_state[sst_support->sup_asid - 1], sst_support);

    // init the term process
    term_pcb[sst_support->sup_asid - 1] = initTermProcess(&term_state[sst_support->sup_asid - 1], sst_support);

    // get the message from someone - user process
    // handle
    // reply
    while (TRUE) {
        ssi_payload_PTR process_request_payload;
        pcb_PTR process_request_ptr = (pcb_PTR)SYSCALL(RECEIVEMESSAGE, ANYMESSAGE, (unsigned)(&process_request_payload), 0);
        sstRequestHandler(process_request_ptr, process_request_payload->service_code,
                          process_request_payload->arg,
                          print_pcb[sst_support->sup_asid - 1],
                          term_pcb[sst_support->sup_asid - 1]);
    }
}

```

The SST services are:

- **GetTOD:** this service should allow the sender to get back the number of microseconds since the system was last booted/reset.

```

cpu_t getTOD() {
    cpu_t tod_time;
    STCK(tod_time);
    return tod_time;
}

```

- **Terminate:** this service causes the sender U-proc and its SST (its parent) to cease to exist. It is essentially a SST “wrapper” for the SSI service **TerminateProcess**. It also marks all of the frames in the swap pool from *occupied* to **unoccupied**: this is accomplished by atomically setting the swap pool entry **asid** to **NOPROC(-1)**.

```

void killSST(pcb_PTR sender) {
    notify(test_process);

    // invalidate the page table
    invalidateUProcPageTable(sst_pcb[asid]->p_supportStruct);

    // kill the sst and its child
    terminateProcess(SELF);
}

```

- **WritePrinter:** this service should allow the sender to write a string to the printer. The string is passed as a parameter in the message. The SST should then send a message to the print process with the string to be printed.
- **WriteTerminal:** the same as the **WritePrinter** but for the terminal.

```

void print(unsigned code, sst_print_PTR arg, pcb_PTR print_process) {
    // unwrap the arg and send it to the print process
    int length = arg->length;
    char string[length];
    for (int i = 0; i < length; i++) {
        string[i] = arg->string[i];
    }
    sst_print_t printing = {
        .string = string,
        .length = length,
    };
    SYSCALL(SENDMESSAGE, (unsigned int)print_process, (unsigned int)&printing, 0);
    SYSCALL(RECEIVEMESSAGE, (unsigned)print_process, 0, 0);
}

```

Stdlib

this file serves as the main container of useful functions that are used through phase 3 of the project such as:

- Initializations functions: *initUprocPageTable()*, *initFreeStackTop()*, *initUProc()* and *defaultSupportData()*, *initHelper()*;
- Utility functions: *getSupportData()*, *getCurrentFreeStackTop()*, *createChild()*, *terminateProcess()*, *updateTLB()*, *invalidateUProcPageTable()*;
- Notification service & Mutual exclusion handling: *notify()*, *gainSwapMutex()* and *releaseSwapMutex()*.
- I/O processes: *initPrintProcess()*, *initTermProcess()*, *termEntry()*, *printEntry()*, *writeOnPrinter()*, *write()*, *writeOnTerminal()*.
- support handling: *deallocateSupport()*, *allocateSupport()*.

I/O processes

The I/O processes are used to handle I/O operations such as writing to the printer and terminal. The I/O processes are implemented as separate processes that are responsible for handling I/O requests from the SST. The I/O processes are initialized by the *initPrintProcess* and *initTermProcess* functions, which create the I/O processes and initialize their state. The I/O processes enter a loop where they wait for messages from the SST, handle the request and send back the result.

The *initHelper* function is a generic way for initialize a machine-mode processes.


```

pcb_PTR initPrintProcess(state_t *print_state, support_t *sst_support) {
    return initHelper(print_state, sst_support, printEntry);
}

pcb_PTR initHelper(state_t *helper_state, support_t *sst_support, void *entry) {
    STST(helper_state);
    helper_state->entry_hi = sst_support->sup_asid << ASIDSHIFT;
    helper_state->pc_epc = (memaddr)entry;
    helper_state->reg_sp = getCurrentFreeStackTop();
    helper_state->status = MSTATUS_MPIE_MASK | MSTATUS_MPP_M | MSTATUS_MIE_MASK;
    helper_state->mie = MIE_ALL;

    return createChild(helper_state, sst_support);
}

void printEntry() {
    support_t *support = getSupportData();
    unsigned asid = support->sup_asid;

    while (TRUE) {
        sst_print_PTR print_payload;
        pcb_PTR sender = (pcb_PTR)SYSCALL(RECEIVEMESSAGE, (unsigned)sst_pcb[asid - 1], (unsigned int)(amp;print_payload), 0);

        writeOnPrinter(print_payload, asid);

        // notify the sender that the print is done
        SYSCALL(SENDMESSAGE, (unsigned)sender, 0, 0);
    }
}

void writeOnPrinter(sst_print_PTR arg, unsigned asid) {
    write(arg->string, arg->length, (devreg_t *)DEV_REG_ADDR(IL_PRINTER, asid - 1), PRINTER);
}

void write(char *msg, int lenght, devreg_t *devAddrBase, enum writet write_to) {
    int i = 0;
    unsigned status;
    // check if it's a terminal or a printer
    unsigned *command = write_to == TERMINAL ? amp;(devAddrBase->term.transm_command)
        : amp;(devAddrBase->dtp.command);

    while (TRUE) {
        if ((*msg == EOS) || (i >= lenght)) {
            break;
        }

        unsigned int value;
        status = 0;

        if (write_to == TERMINAL) {
            value = PRINTCHR | (((unsigned int)*msg) << 8);
        } else {
            value = PRINTCHR;
            devAddrBase->dtp.data0 = *msg;
        }

        ssi_do_io_t do_io = {
            .commandAddr = command,
            .commandValue = value,

```

```

};

ssi_payload_t payload = {
    .service_code = DOIO,
    .arg = &do_io,
};

SYSCALL(SENDMESSAGE, (unsigned int)ssi_pcb, (unsigned int>(&payload), 0);
SYSCALL(RECEIVEMESSAGE, (unsigned int)ssi_pcb, (unsigned int>(&status), 0);

// device not ready -> error!
if (write_to == TERMINAL && status != OKCHARTRANS) {
    terminateParent();
} else if (write_to == PRINTER && status != DEVRDY) {
    terminateParent();
}

msg++;
i++;
}
}

```

Note that the same logic is applied to the terminal process.

Support handling

The support handling is a set of functions that are used to manage the support struct of the processes. The support struct is a data structure that is used to store the state of the process, such as the page table, the ASID, and the TLB entries. The support handling functions are used to allocate and deallocate the support struct, and to get the support struct of the current process. The way it's accomplished is by using a global array of support structs then a free support struct list. The `allocateSupport` function is used to allocate a support struct from the free list, and the `deallocateSupport` function is used to deallocate a support struct and add it back to the free list.

```

support_t *allocateSupport(void) {
    static int asid = 1;
    if (list_empty(&free_supports) || asid > MAXSSTNUM) {
        return NULL;
    }

    struct list_head *head = free_supports.next;
    list_del(head);
    support_t *s = container_of(head, support_t, s_list);
    s->sup_asid = asid++;

    defaultSupportData(s, s->sup_asid);
    return s;
}

void deallocateSupport(support_t *s) {
    if (s->sup_asid != 0) {
        invalidateUProcPageTable(s->sup_asid);
    }

    list_add(&s->s_list, &free_supports);
}

void defaultSupportData(support_t *support_data, int asid) {
    support_data->sup_asid = asid;

    support_data->sup_exceptContext[PGFAULTEXCEPT].pc = (memaddr)pager;
    support_data->sup_exceptContext[PGFAULTEXCEPT].stackPtr = getCurrentFreeStackTop();
    support_data->sup_exceptContext[PGFAULTEXCEPT].status = MSTATUS_MIE_MASK | MSTATUS_MPP_M | MSTATUS_MPIE_MASK;

    support_data->sup_exceptContext[GENERALEXCEPT].pc = (memaddr)supportExceptionHandler;
    support_data->sup_exceptContext[GENERALEXCEPT].stackPtr = getCurrentFreeStackTop();
    support_data->sup_exceptContext[GENERALEXCEPT].status = MSTATUS_MIE_MASK | MSTATUS_MPIE_MASK | MSTATUS_MPP_M;

    initUprocPageTable(support_data->sup_privatePgTbl, asid);

    INIT_LIST_HEAD(&support_data->s_list);
}

```

P3test

This file serves as the main test for the phase 3, it follow a simple schema: initialization, execution of the process's request and termination (after being notified).

```
void test3() {
    test_process = current_process;
    initFreeStackTop();

    // Init array of support struct (so each will be used for every u-proc init. in initSSTs)
    initSupportArray();

    // alloc swap mutex process
    swap_mutex = allocSwapMutex();

    //Init 8 SST
    initSSTs();

    //Terminate after the 8 sst die
    waitTermination(sst_pcb);

    // terminate the test process
    terminateProcess(SELF);
}
```

It also contains some utility functions such as: *initSupportArray()*, *allocSwapMutex()* and *waitTermination()*.